

hausaTweet

July 3, 2021

```
[1]: #import libraries

import pandas as pd
import numpy as np
import sklearn.metrics as metrics
import matplotlib.pyplot as plt
import seaborn as sns
import nltk

#import Skit-helper functions
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer

#import Scikit-learn models
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression

#import performance metrics functions
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, \
    recall_score
from sklearn.metrics import classification_report
from sklearn.metrics import auc, roc_auc_score, roc_curve
```

```
[2]: #importing dataset

data = pd.read_csv('sentiment_analysis_ouptut.csv')
data
```

```
[2]:
```

	linenumber	sentiment	tweet_text
0	1	2	Wllh ana biya sbd kudin beli yawa ne dasu har ...
1	2	3	Kowa yayi da damisa ya duba jikinsa.
2	3	3	Duk wanda ka gani a cikin motar kai shi ka yi ...
3	4	-1	Allah ya tsinewa masu 'karya
4	5	-2	Mutane Basu D Adalcin Magana, Shugaba Yace Zai...
...
2944	4193	1	Muma burin haka. Domain wanna makircine ga zam...
2945	4194	-2	Ay dama raini ne da rashin hankali.. Amma in b...

2946	4196	-1	Shikenan Wanda aka konawa ababen hawa sai dai ...
2947	4198	7	Wannan abun a yaba ne, domin hakan zai sake in...
2948	4199	-1	Na zata rivers bata cikin IPOB

[2949 rows x 3 columns]

```
[3]: #Categorizing sentiments
#(-1 downward as negative sentiment category[0])
#(+1 upward as positive sentiment category [1])

data.loc[(data['sentiment'] <0 ), 'sentiment'] = 0
data.loc[(data['sentiment'] >0), 'sentiment'] = 1

j = 0
for i in range(0,1):
    data.loc[(data["sentiment"] >= j) & (data["sentiment"] <= i*10),
    ↪ "sentiment"] = i*10
    i = i + 1
    j = j + 10

data
```

```
[3]:      linenumber  sentiment      tweet_text
0           1          1  Wllh ana biya sbd kudin beli yawa ne dasu har ...
1           2          1           Kowa yayi da damisa ya duba jikinsa.
2           3          1  Duk wanda ka gani a cikin motar kai shi ka yi ...
3           4          0           Allah ya tsinewa masu 'karya
4           5          0  Mutane Basu D Adalcin Magana, Shugaba Yace Zai...
...         ...         ...
2944        4193          1  Muma burin haka. Domain wanna makircine ga zam...
2945        4194          0  Ay dama raini ne da rashin hankali.. Amma in b...
2946        4196          0  Shikenan Wanda aka konawa ababen hawa sai dai ...
2947        4198          1  Wannan abun a yaba ne, domin hakan zai sake in...
2948        4199          0           Na zata rivers bata cikin IPOB
```

[2949 rows x 3 columns]

```
[4]: #data split into training and testing dataset

test_percentage = .2
train_df, test_df = train_test_split(data, test_size=test_percentage,
    ↪ random_state=42)

labels = train_df['sentiment']
test_labels = test_df['sentiment']

print("\n### Split Complete ###\n")
```

Split Complete

```
[5]: # Print counts of each class
print("- Counting Splits -")
print("Training Samples:", len(train_df))
print("Testing Samples:", len(test_df))

# Graph counts of each class, for both training and testing


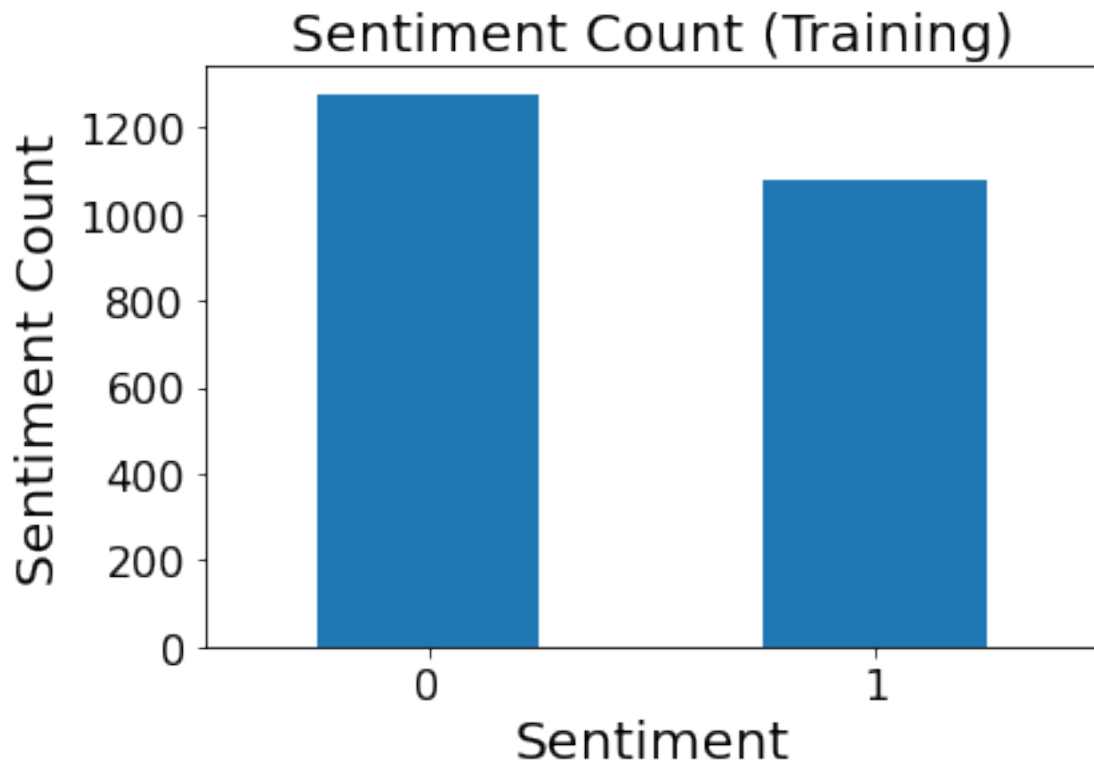
count_train_classes = pd.value_counts(train_df['sentiment'])
count_train_classes.plot(kind='bar', fontsize=16)
plt.title("Sentiment Count (Training)", fontsize=20)
plt.xticks(rotation='horizontal')
plt.xlabel("Sentiment", fontsize=20)
plt.ylabel("Sentiment Count", fontsize=20)

plt.show()

count_test_classes = pd.value_counts(test_df['sentiment'])
count_test_classes.plot(kind='bar', fontsize=16, colormap='ocean')
plt.title("Sentiment Count (Testing)", fontsize=20)
plt.xticks(rotation='horizontal')
plt.xlabel("Sentiment", fontsize=20)
plt.ylabel("Sentiment Count", fontsize=20)

plt.show()
```

- Counting Splits -
Training Samples: 2359
Testing Samples: 590



```
[6]: # Vectorizer the training inputs -- Takes about 30 seconds to complete
#     There are two types of vectors:
#         1. Count vectorizer
#         2. Term Frequency-Inverse Document Frequency (TF-IDF)
from nltk.tokenize import word_tokenize

print("- Training Count Vectorizer -")
cVec = CountVectorizer()
count_X = cVec.fit_transform(train_df['tweet_text'])

print("- Training TF-IDF Vectorizer -")
tVec = TfidfVectorizer()
tfidf_X = tVec.fit_transform(train_df['tweet_text'])

print("\n### Vectorizing Complete ###\n")
```

```
- Training Count Vectorizer -
- Training TF-IDF Vectorizer -
```

```
### Vectorizing Complete ###
```

```
[7]: # Vectorize the testing inputs
#     Use 'transform' instead of 'fit_transform' because we've already trained
#     ↳ our vectorizers

print("- Count Vectorizer -")
test_count_X = cVec.transform(test_df['tweet_text'])

print("- TFIDF Vectorizer -")
test_tfidf_X = tVec.transform(test_df['tweet_text'])

print("\n### Vectorizing Complete ###\n")
```

```
- Count Vectorizer -
- TFIDF Vectorizer -
```

```
### Vectorizing Complete ###
```

```
[8]: def generate_report(cmatrix, score, creport):
      """Generates and displays graphical reports
```

```

Keyword arguments:
    cmatrix - Confusion matrix generated by the model
    score --- Score generated by the model
    creport - Classification Report generated by the model

:Returns -- N/A
"""

# Transform cmatrix because Sklearn has pred as columns and actual as rows.
cmatrix = cmatrix.T

# Generate confusion matrix heatmap
plt.figure(figsize=(5,5))
sns.heatmap(cmatrix,
            annot=True,
            fmt="d",
            linewidths=.5,
            square = True,
            cmap = 'Blues',
            annot_kws={"size": 16},
            xticklabels=['negative', 'positive'],
            yticklabels=['negative', 'positive'])

plt.xticks(rotation='horizontal', fontsize=16)
plt.yticks(rotation='horizontal', fontsize=16)
plt.xlabel('Actual Label', size=20);
plt.ylabel('Predicted Label', size=20);

title = 'Accuracy Score: {0:.4f}'.format(score)
plt.title(title, size = 20);

# Display classification report and confusion matrix
print(creport)
plt.show()

print("\n### Report Generator Defined ###\n")

```

```
### Report Generator Defined ###
```

```
[9]: # Multinomial Naive Bayesian with TF-IDF
```

```

# Train the model
mnb_tfidf = MultinomialNB()
mnb_tfidf.fit(tfidf_X, labels)

```

```


# Test the mode (score, predictions, confusion matrix, classification report)
score_mnb_tfidf = mnb_tfidf.score(test_tfidf_X, test_labels)
predictions_mnb_tfidf = mnb_tfidf.predict(test_tfidf_X)
cmatrix_mnb_tfidf = confusion_matrix(test_labels, predictions_mnb_tfidf)
creport_mnb_tfidf = classification_report(test_labels, predictions_mnb_tfidf)

print("\n### Model Built ###\n")
generate_report(cmatrix_mnb_tfidf, score_mnb_tfidf, creport_mnb_tfidf)

```

Model Built

	precision	recall	f1-score	support
0	0.77	0.87	0.82	306
1	0.84	0.72	0.77	284
accuracy			0.80	590
macro avg	0.80	0.80	0.80	590
weighted avg	0.80	0.80	0.80	590



```
[10]: # Multinomial Naive Bayesian with Count Vectorizer


# Train the model
mnb_count = MultinomialNB()
mnb_count.fit(count_X, labels)

# Test the mode (score, predictions, confusion matrix, classification report)
score_mnb_count = mnb_count.score(test_count_X, test_labels)
predictions_mnb_count = mnb_count.predict(test_count_X)
cmatrix_mnb_count = confusion_matrix(test_labels, predictions_mnb_count)
creport_mnb_count = classification_report(test_labels, predictions_mnb_count)

print("\n### Model Built ###\n")
generate_report(cmatrix_mnb_count, score_mnb_count, creport_mnb_count)
```

Model Built

	precision	recall	f1-score	support
0	0.84	0.76	0.80	306
1	0.77	0.84	0.80	284
accuracy			0.80	590
macro avg	0.80	0.80	0.80	590
weighted avg	0.80	0.80	0.80	590



```
[11]: # Logistic Regression with TF-IDF

# Train the model
lgs_tfidf = LogisticRegression(solver='lbfgs')
lgs_tfidf.fit(tfidf_X, labels)


# Test the mode (score, predictions, confusion matrix, classification report)
score_lgs_tfidf = lgs_tfidf.score(test_tfidf_X, test_labels)
predictions_lgs_tfidf = lgs_tfidf.predict(test_tfidf_X)
cmatrix_lgs_tfidf = confusion_matrix(test_labels, predictions_lgs_tfidf)
creport_lgs_tfidf = classification_report(test_labels, predictions_lgs_tfidf)

print("\n### Model Built ###\n")
generate_report(cmatrix_lgs_tfidf, score_lgs_tfidf, creport_lgs_tfidf)
```

Model Built

	precision	recall	f1-score	support
0	0.79	0.93	0.85	306

	1	0.90	0.73	0.81	284
accuracy				0.83	590
macro avg		0.85	0.83	0.83	590
weighted avg		0.84	0.83	0.83	590



```
[12]: # Logistic Regression with Count Vectorizer

# Train the model
lgs_count = LogisticRegression(solver='lbfgs')
lgs_count.fit(count_X, labels)

# Test the mode (score, predictions, confusion matrix, classification report)
score_lgs_count = lgs_count.score(test_count_X, test_labels)
predictions_lgs_count = lgs_count.predict(test_count_X)
cmatrix_lgs_count = confusion_matrix(test_labels, predictions_lgs_count)
creport_lgs_count = classification_report(test_labels, predictions_lgs_count)

print("\n### Model Built ###\n")
generate_report(cmatrix_lgs_count, score_lgs_count, creport_lgs_count)
```

Model Built

	precision	recall	f1-score	support
0	0.83	0.93	0.88	306
1	0.91	0.80	0.85	284
accuracy			0.86	590
macro avg	0.87	0.86	0.86	590
weighted avg	0.87	0.86	0.86	590

